

HPC Lab Manual: Distributed Computing and Parallel Programming

Table of Contents

1. [System Access](#)
 2. [Directory Structure Overview](#)
 3. [SLURM Job Scheduler](#)
 4. [Distributed Machine Learning](#)
 5. [Parallel Programming with C/C++](#)
 6. [Common Commands Reference](#)
 7. [Troubleshooting](#)
-

System Access

SSH Connection

```
bash  
ssh username@master-node
```

Replace `username` with your actual username.

Navigate to Summer School Directory

```
bash  
cd /mnt/lustre/nabeel/summer_school
```

Directory Listing

```
bash  
ls -la
```

Directory Structure Overview

The summer school materials are organized in the following hierarchy:

/mnt/lustre/nabeel/summer_school/

```
|— Demonstration_Mnist_Dist/      # Data Level Parallelism (MNIST 70k)
|   |— data/                      # Dataset files
|   |— nccl_1259.err              # Error logs
|   |— nccl_1259.out              # Output logs
|   |— py_train.slurm             # Python training SLURM script
|   |— train.py                   # Training script
|— Dist_DL/                       # Distributed Deep Learning
|   |— GPU_Off/                  # GPU Off examples
|   |— model_shard/              # Model Level Parallelism (Model Sharding)
|— job_sch_new/                   # Job Scheduling & Distributed Training (10k dataset)
|   |— comm_1264.err              # Communication error logs
|   |— dataset.py                # Dataset handling
|   |— ft_1263.err               # Fault tolerance error logs
|   |— ft_1263.out               # Fault tolerance output logs
|   |— job_sched1.slurm          # Main scheduling script
|   |— model.py                  # Model definition
|   |— nccl_1261.err             # NCCL error logs
|   |— nccl_1261.out             # NCCL output logs
|   |— nccl_1262.err             # Additional NCCL error logs
|   |— nccl_1262.out             # Additional NCCL output logs
|   |— __pycache__/              # Python cache
|   |— run_local.sh              # Local execution script
|   |— slurm_comm.sh             # Node communication script
|   |— slurm_ft.slurm            # Fault tolerance SLURM script
|   |— train.py                  # Training script with time metrics
|— Job_Sch_Res_Man/              # Job Scheduling Resource Management
|— MPI/                          # Message Passing Interface
|— openmp_prog/                  # OpenMP Programming
|— SMP/                          # Symmetric Multi-Processing
```

SLURM Job Scheduler

Basic SLURM Script Structure

Create a file named `job_script.slurm`:

```
bash
```

```
#!/bin/bash

#SBATCH --job-name=my_job      # Job name
#SBATCH --nodes=1             # Number of nodes
#SBATCH --ntasks-per-node=4   # Tasks per node
#SBATCH --cpus-per-task=4     # CPUs per task
#SBATCH --gpus-per-node=1     # GPUs per node (if needed)
#SBATCH --time=01:00:00       # Time limit (HH:MM:SS)
#SBATCH --output=nccl_%j.out   # Output file
#SBATCH --error=nccl_%j.err    # Error file


# Master node address and port
export MASTER_ADDR=192.168.20.15
export MASTER_PORT=29500


# Set Ethernet interface for NCCL
export NCCL_SOCKET_IFNAME=ensif1


# Disable InfiniBand (if not available)
export NCCL_IB_DISABLE=1


# Optional: Enable debug output
export NCCL_DEBUG=INFO
export PYTHONUNBUFFERED=1


# Export distributed environment variables
export RANK
export WORLD_SIZE
export MASTER_ADDR
export MASTER_PORT


echo "RANK=$RANK on $(hostname) using interface $NCCL_SOCKET_IFNAME"


# Run your program
python3 train.py
```

SLURM Parameter Explanation

Parameter	Description	Example Values
<code>--job-name</code>	Name for your job	<code>mnist_training</code> , <code>mpi_test</code>
<code>--nodes</code>	Number of compute nodes	<code>1</code> , <code>2</code> , <code>4</code>
<code>--ntasks-per-node</code>	Tasks per node	<code>1</code> , <code>4</code> , <code>8</code>
<code>--cpus-per-task</code>	CPU cores per task	<code>1</code> , <code>4</code> , <code>8</code>
<code>--gpus-per-node</code>	GPUs per node	<code>1</code> , <code>2</code> , <code>4</code>
<code>--time</code>	Maximum runtime	<code>00:30:00</code> , <code>02:00:00</code>
<code>--partition</code>	Queue/partition name	<code>gpu</code> , <code>cpu</code> , <code>debug</code>
<code>--mem</code>	Memory per node	<code>8G</code> , <code>16G</code> , <code>32G</code>

Essential SLURM Commands

Submit a Job

```
bash

sbatch job_script.slurm
```

Check Job Status

```
bash

squeue           # All jobs
squeue -u username  # Your jobs only
squeue -j jobid    # Specific job
```

Cancel a Job

```
bash

scancel jobid      # Cancel specific job
scancel -u username  # Cancel all your jobs
```

Job Information

```
bash

sinfo           # Node information
sacct -j jobid   # Job accounting info
```

Distributed Machine Learning

Data Level Parallelism (MNIST 70,000 Images)

Navigate to the MNIST distributed training directory:

```
bash  
  
cd /mnt/lustre/nabeel/summer_school/Demonstration_Mnist_Dist
```

Directory Contents:

- `data/` - Contains the MNIST dataset
- `train.py` - Main training script for distributed learning
- `py_train.slurm` - SLURM batch script for Python training
- `nccl_*.out` - Training output logs
- `nccl_*.err` - Error logs

Running MNIST Distributed Training:

Step 1: Review the SLURM Script

```
bash  
  
cat py_train.slurm
```

Step 2: Submit the Job

```
bash  
  
sbatch py_train.slurm
```

Step 3: Monitor Job Progress

```
bash  
  
# Check job status  
squeue -u $USER  
  
# Monitor real-time output  
tail -f nccl_*.out  
  
# Check for errors  
tail -f nccl_*.err
```

Step 4: View Results

```
bash
```

```
# View complete output
```

```
cat nccl_*.out
```

```
# Check training metrics and performance
```

```
grep -i "epoch\\|loss\\|accuracy" nccl_*.out
```

Key Features:

- Uses 70,000 MNIST images for training
- Implements data parallelism across multiple nodes
- NCCL backend for efficient communication
- Automatic logging of training progress

Distributed Deep Learning Options

GPU Off Examples

Navigate to the GPU off directory:

```
bash
```

```
cd /mnt/lustre/nabeel/summer_school/Dist_DL/GPU_Off
```

Model Level Parallelism (Model Sharding)

Navigate to the model sharding directory:

```
bash
```

```
cd /mnt/lustre/nabeel/summer_school/Dist_DL/model_shard
```

Running Model Sharding

```
bash
```

```
# Submit the job
```

```
sbatch model_shard.slurm
```

```
# Check logs
```

```
cat model_shard_*.out
```

Job Scheduling and Time Metrics (10,000 Image Dataset)

Navigate to the job scheduling directory:

```
bash
```

```
cd /mnt/lustre/nabeel/summer_school/job_sch_new
```

Directory Contents:

- `train.py` - Training script with comprehensive time metrics
- `model.py` - Neural network model definition
- `dataset.py` - Dataset handling and preprocessing
- `job_sched1.slurm` - **Main scheduling script for training**
- `slurm_ft.slurm` - **Fault tolerance SLURM script**
- `slurm_comm.sh` - **Node communication script**
- `run_local.sh` - Local execution script
- `nccl_*.out/err` - Output and error logs
- `ft_*.out/err` - Fault tolerance logs
- `comm_*.err` - Communication logs

Running Job Scheduling Experiments:

Option 1: Standard Distributed Training with Scheduling

```
bash
```

```
# Submit the main scheduling job
```

```
sbatch job_sched1.slurm
```

```
# Monitor progress
```

```
squeue -u $USER
```

```
tail -f nccl_*.out
```

Option 2: Fault Tolerance Training

```
bash
```

```
# Submit fault tolerance job
```

```
sbatch slurm_ft.slurm
```

```
# Monitor fault tolerance logs
```

```
tail -f ft_*.out
```

Option 3: Node Communication Testing

```
bash
```

```
# Run communication script
```

```
bash slurm_comm.sh
```

```
# Check communication logs
```

```
tail -f comm_*.err
```

Option 4: Local Testing

```
bash
```

```
# Run locally for debugging
```

```
bash run_local.sh
```

Changing Number of Nodes for Performance Testing:

Edit SLURM Scripts to Test Different Node Configurations:

```
bash
```

```
# Edit the main scheduling script
```

```
nano job_sched1.slurm
```

```
# Modify these parameters:
```

```
#SBATCH --nodes=1      # Change to 1, 2, 4, 8 nodes
```

```
#SBATCH --ntasks-per-node=4 # Adjust tasks per node
```

Example configurations for testing:

Single Node:

```
bash
```

```
#SBATCH --nodes=1
```

```
#SBATCH --ntasks-per-node=4
```

Two Nodes:

```
bash
```

```
#SBATCH --nodes=2
```

```
#SBATCH --ntasks-per-node=4
```

Four Nodes:

```
bash
```



```
#SBATCH --nodes=4
#SBATCH --ntasks-per-node=2
```

Analyzing Time Metrics:

View Training Time Results:

```
bash

# Check training completion times
grep -i "training time\|epoch time\|total time" nccl_*.out

# Compare performance across different node configurations
grep -i "throughput\|images/sec\|batch time" nccl_*.out
```

Performance Comparison Workflow:

1. Submit job with 1 node: `sbatch job_sched1.slurm`
2. Record job ID and wait for completion
3. Edit script to use 2 nodes, submit again
4. Compare time metrics in output files
5. Repeat for 4, 8 nodes to analyze scaling

Key Features:

- Uses 10,000 image dataset for faster experimentation
- Comprehensive time metrics and performance monitoring
- Multiple execution modes (standard, fault tolerance, communication testing)
- Easy node configuration changes for scaling experiments
- Detailed logging for performance analysis

Job Scheduling Resource Management

Navigate to the resource management directory:

```
bash

cd /mnt/lustre/nabeel/summer_school/Job_Sch_Res_Man
```

This directory focuses on:

- Resource allocation strategies
- Job queue management

- System resource optimization
 - Scheduling algorithms
-

Parallel Programming with C/C++

MPI (Message Passing Interface)

```
bash
cd /mnt/lustre/nabeel/summer_school/MPI
```

Compilation:

```
bash
mpicc -o program program.c      # C program
mpicxx -o program program.cpp   # C++ program
```

SLURM Script for MPI:

```
bash
#!/bin/bash
#SBATCH --job-name=mpi_job
#SBATCH --nodes=2
#SBATCH --ntasks=8
#SBATCH --time=00:30:00

mpirun -np 8 ./program
```

OpenMP Programming

```
bash
cd /mnt/lustre/nabeel/summer_school/openmp_prog
```

Compilation:

```
bash
gcc -fopenmp -o program program.c # GCC
icc -qopenmp -o program program.c # Intel compiler
```

SLURM Script for OpenMP:

```
bash
```

```
#!/bin/bash
```

```
#SBATCH --job-name=openmp_job
```

```
#SBATCH --nodes=1
```

```
#SBATCH --cpus-per-task=8
```

```
#SBATCH --time=00:30:00
```

```
export OMP_NUM_THREADS=8
```

```
./program
```

Symmetric Multi-Processing (SMP)

```
bash
```

```
cd /mnt/lustre/nabeel/summer_school/SMP
```

Compilation:

```
bash
```

```
gcc -o program program.c      # Basic compilation
```

```
gcc -O3 -o program program.c  # Optimized compilation
```

SLURM Script for SMP:

```
bash
```

```
#!/bin/bash
```

```
#SBATCH --job-name=smp_job
```

```
#SBATCH --nodes=1
```

```
#SBATCH --cpus-per-task=4
```

```
#SBATCH --time=00:30:00
```

```
./program
```

Common Commands Reference

File Operations

```
bash
```

```
ls -la                # List files with details
cp file1 file2        # Copy file
mv file1 file2        # Move/rename file
rm file               # Delete file
mkdir dirname         # Create directory
```

Viewing Files

```
bash

cat filename          # Display entire file
head -n 20 filename   # First 20 lines
tail -n 20 filename   # Last 20 lines
tail -f filename      # Follow file updates
less filename         # Page through file
```

Process Management

```
bash

ps aux                # List all processes
top                   # System monitor
htop                  # Enhanced system monitor
kill -9 PID           # Kill process
```

System Information

```
bash

nvidia-smi            # GPU status
free -h               # Memory usage
df -h                 # Disk usage
lscpu                  # CPU information
```

Troubleshooting

Common Issues and Solutions

Job Fails to Start

1. Check SLURM script syntax
2. Verify resource availability: `sinfo`
3. Check job queue: `squeue`

Out of Memory Errors

```
bash
```

```
#SBATCH --mem=32G           # Increase memory allocation
```

GPU Not Detected

```
bash
```

```
#SBATCH --gpus-per-node=1    # Explicitly request GPU  
export CUDA_VISIBLE_DEVICES=0 # Set GPU device
```

NCCL Communication Issues

```
bash
```

```
export NCCL_DEBUG=INFO        # Enable debug output  
export NCCL_IB_DISABLE=1      # Disable InfiniBand  
export NCCL_SOCKET_IFNAME=ensif1 # Set network interface
```

Python Path Issues

```
bash
```

```
export PYTHONPATH=/path/to/modules:$PYTHONPATH  
which python3          # Verify Python location
```

Log File Locations

Check these files for debugging:

- `nccl_*.out` - Standard output
- `nccl_*.err` - Error messages
- `slurm-*.out` - SLURM job output
- Individual directory logs for specific applications

Getting Help

1. Check log files first
 2. Use `squeue` to verify job status
 3. Review SLURM script parameters
 4. Test on smaller datasets/shorter time limits
 5. Check system resource availability with `sinfo`
-

Quick Start Checklist

1. **Connect to HPC:** `ssh username@master-node`
2. **Navigate to directory:** `cd /mnt/lustre/nabeel/summer_school`
3. **Choose your task:**
 - Data parallelism: `cd Demonstration_Mnist_Dist`
 - Model parallelism: `cd Dist_DL/model_shard`
 - Job scheduling: `cd job_sch_new`
 - MPI/OpenMP/CUDA: `cd MPI/` or `cd OpenMP/` etc.
4. **Review the SLURM script:** `cat job_script.slurm`
5. **Submit job:** `sbatch job_script.slurm`
6. **Monitor:** `squeue -u username`
7. **Check results:** `tail -f output_file.out`

Remember to always check the specific requirements and parameters for each exercise in their respective directories!